

CKA Exam Questions and Answers

Q1) Create a new service account with the name pvviewer. Grant this Service account access to list all PersistentVolumes in the cluster by creating an appropriate cluster role called pvviewer-role and ClusterRoleBinding called pvviewer-role-binding.

Next, create a pod called **pvviewer** with the **image: redis** and **serviceaccount: pvviewer** in the default namespace.

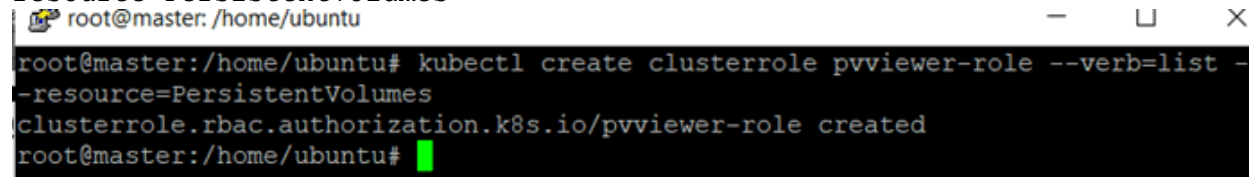
Ans.

Create Service account

```
$ kubectl create serviceaccount pvviewer
```

Create cluster role

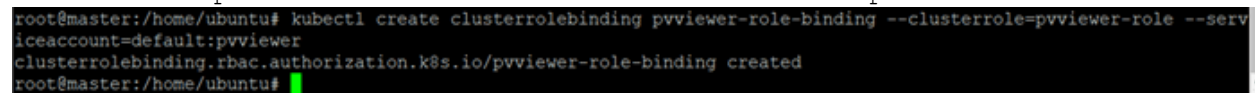
```
$ kubectl create clusterrole pvviewer-role --verb=list --resource=PersistentVolumes
```



```
root@master:/home/ubuntu# kubectl create clusterrole pvviewer-role --verb=list --resource=PersistentVolumes
clusterrole.rbac.authorization.k8s.io/pvviewer-role created
root@master:/home/ubuntu#
```

Create cluster role binding

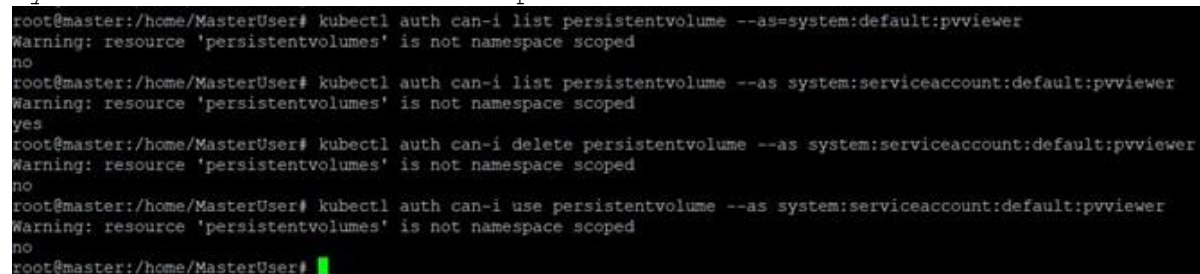
```
$ kubectl create clusterrolebinding pvviewer-role-binding --clusterrole=pvviewer-role --serviceaccount=default:pvviewer
```



```
root@master:/home/ubuntu# kubectl create clusterrolebinding pvviewer-role-binding --clusterrole=pvviewer-role --serviceaccount=default:pvviewer
clusterrolebinding.rbac.authorization.k8s.io/pvviewer-role-binding created
root@master:/home/ubuntu#
```

Verify

```
$ kubectl auth can-i list PersistentVolumes --as system:serviceaccount:default:pvviewer
```



```
root@master:/home/MasterUser# kubectl auth can-i list persistentvolume --as=system:default:pvviewer
Warning: resource 'persistentvolumes' is not namespace scoped
no
root@master:/home/MasterUser# kubectl auth can-i list persistentvolume --as system:serviceaccount:default:pvviewer
Warning: resource 'persistentvolumes' is not namespace scoped
yes
root@master:/home/MasterUser# kubectl auth can-i delete persistentvolume --as system:serviceaccount:default:pvviewer
Warning: resource 'persistentvolumes' is not namespace scoped
no
root@master:/home/MasterUser# kubectl auth can-i use persistentvolume --as system:serviceaccount:default:pvviewer
Warning: resource 'persistentvolumes' is not namespace scoped
no
root@master:/home/MasterUser#
```

Q2) Create a new deployment called nginx-deploy, with image nginx:1.16 and 1 replica. Record the version. Next upgrade the deployment to version 1.17 using rolling update. Make sure that the version upgrade is recorded in the resource annotation.

Ans.

```

apiVersion: apps/v1
kind: Deployment
metadata:
  name: nginx-deploy
  labels:
    app: nginx
spec:
  replicas: 1
  selector:
    matchLabels:
      app: nginx
  template:
    metadata:
      labels:
        app: nginx
    spec:
      containers:
      - name: nginx
        image: nginx:1.16
        ports:
        - containerPort: 80
~
~
~
~
~

```

```

$ vim nginx-deployment.yaml
$ kubectl apply -f nginx-deployment.yaml --record
$ kubectl get deployment
$ kubectl rollout history deployment nginx-deploy
root@kubeadm-master:/home/ubuntu/Kubernetes# kubectl get deployment
NAME          READY   UP-TO-DATE   AVAILABLE   AGE
nginx-deploy   1/1     1            1           2m22s
root@kubeadm-master:/home/ubuntu/Kubernetes# kubectl rollout history deployment nginx-deploy
deployment.apps/nginx-deploy
REVISION  CHANGE-CAUSE
1          kubectl apply --filename=nginx-deployment.yaml --record=true

root@kubeadm-master:/home/ubuntu/Kubernetes# █

```

```

$ kubectl set image deployment/nginx-deploy nginx=1.17 --record
$ kubectl rollout history deployment nginx-deploy
root@kubeadm-master:/home/ubuntu/Kubernetes# kubectl set image deployment/nginx-deploy nginx=1.17 --record
deployment.apps/nginx-deploy image updated
root@kubeadm-master:/home/ubuntu/Kubernetes# kubectl rollout history deployment nginx-deploy
deployment.apps/nginx-deploy
REVISION  CHANGE-CAUSE
1          kubectl apply --filename=nginx-deployment.yaml --record=true
2          kubectl set image deployment/nginx-deploy nginx=1.17 --record=true

root@kubeadm-master:/home/ubuntu/Kubernetes# █

```

```

$ kubectl describe deployment nginx-deploy

```

```

root@kubeadm-master:/home/ubuntu/Kubernetes# kubectl describe deployment nginx-deploy
Name:          nginx-deploy
Namespace:     default
CreationTimestamp: Mon, 21 Sep 2020 05:34:39 +0000
Labels:        app=nginx
Annotations:    deployment.kubernetes.io/revision: 2
                kubernetes.io/change-cause: kubectl set image deployment/nginx-deploy nginx=1.17 --record=true
Selector:      app=nginx
Replicas:      1 desired | 1 updated | 2 total | 1 available | 1 unavailable
StrategyType:  RollingUpdate
MinReadySeconds: 0
RollingUpdateStrategy: 25% max unavailable, 25% max surge
Pod Template:
  Labels:  app=nginx
  Containers:
    nginx:
      Image:      1.17
      Port:       80/TCP
      Host Port:  0/TCP
      Environment: <none>
      Mounts:      <none>
      Volumes:     <none>
  Conditions:
    Type           Status    Reason
    ----           -
    Available      True      MinimumReplicasAvailable
    Progressing    True      ReplicaSetUpdated
  OldReplicaSets:  nginx-deploy-767cbb69b8 (1/1 replicas created)
  NewReplicaSet:   nginx-deploy-649f54f665 (1/1 replicas created)
Events:
  Type           Reason              Age           From              Message
  ----           -
  Normal         ScalingReplicaSet   3m14s        deployment-controller Scaled up replica set nginx-deploy-767cbb69b8 to 1
  Normal         ScalingReplicaSet   30s          deployment-controller Scaled up replica set nginx-deploy-649f54f665 to 1
root@kubeadm-master:/home/ubuntu/Kubernetes#

```

Q3) Create snapshot of the etcd running at <https://127.0.0.1:2379>. Save snapshot into /opt/etcd-snapshot.db.

Use these are certificate for snapshot

Ca certificate: /etc/kubernetes/pki/etcd/ca.crt
 Client certificate: /etc/kubernetes/pki/etcd/server.crt
 client key: /etc/kubernetes/pki/etcd/server.key
 and then restore from the previous ETCD backup.

Ans:

```

$ ETCDCTL_API=3 etcdctl --endpoints=https://127.0.0.1:2379 --
cert=/etc/kubernetes/pki/etcd/server.crt --
cacert=/etc/kubernetes/pki/etcd/ca.crt --
key=/etc/kubernetes/pki/etcd/server.key snapshot save /opt/etcd-
snapshot.db
Verify//

```

Note: Do not perform this step in exam otherwise it may create an issue in the restoration process.

```

$ ETCDCTL_API=3 etcdctl --write-out=table snapshot status /opt/etcd-
snapshot.db

```

Restore

No need to remember all the flags in the restore command:

You can do

```
$ ETCDCTL_API=3 etcdctl snapshot restore -h
```

```
root@master:/home/ubuntu# ETCDCTL_API=3 etcdctl snapshot restore -h
NAME:
    snapshot restore - Restores an etcd member snapshot to an etcd directory

USAGE:
    etcdctl snapshot restore <filename> [options] [flags]

OPTIONS:
    --data-dir=""                Path to the data directory
    -h, --help[=false]           help for restore
    --initial-advertise-peer-urls="http://localhost:2380" List of this member's peer URLs to advertise to the rest of the cluster
    --initial-cluster="default=http://localhost:2380"    Initial cluster configuration for restore bootstrap
    --initial-cluster-token="etcd-cluster"              Initial cluster token for the etcd cluster during restore bootstrap
    --name="default"                                    Human-readable name for this member
    --skip-hash-check[=false]                           Ignore snapshot integrity hash value (required if copied from data directory)

GLOBAL OPTIONS:
    --cacert=""                verify certificates of TLS-enabled secure servers using this CA bundle
    --cert=""                  identify secure client using this TLS certificate file
    --command-timeout=5s       timeout for short running command (excluding dial timeout)
    --debug[=false]           enable client-side debug logging
    --dial-timeout=2s          dial timeout for client connections
    --endpoints={127.0.0.1:2379} gRPC endpoints
    --hex[=false]              print byte strings as hex encoded strings
    --insecure-skip-tls-verify[=false] skip server certificate verification
    --insecure-transport[=true] disable transport security for client connections
    --key=""                   identify secure client using this TLS key file
    --user=""                  username[:password] for authentication (prompt if password is not supplied)
    -w, --write-out="simple"    set the output format (fields, json, protobuf, simple, table)

root@master:/home/ubuntu#
```

```
$ ETCDCTL_API=3 etcdctl snapshot restore /opt/etcd-snapshot.db --
endpoints=https://127.0.0.1:2379 --
cert=/etc/kubernetes/pki/etcd/server.crt --
cacert=/etc/kubernetes/pki/etcd/ca.crt --
key=/etc/kubernetes/pki/etcd/server.key --data-dir=/var/lib/etcd --
initial-advertise-peer-urls=http://10.0.0.4:2380 --initial-
cluster=<master-name>=http://10.0.0.4:2380" --initial-cluster-token="etcd-
cluster" --name="<master-name>"
```

Q4) Create a Persistent Volume with the given specification.

Volume Name: pv-analytics, Storage: 100Mi, Access modes: ReadWriteMany, Host Path: /pv/data-analytics

Ans.

```
$ vim pv.yaml
apiVersion: v1
kind: PersistentVolume
metadata:
  name: pv-analytics
spec:
  capacity:
    storage: 100Mi
  accessModes:
    - ReadWriteMany
  hostPath:
    path: /pv/data-analytics
$ kubectl create -f pv.yaml
$ kubectl get pv
```

```
root@master:~# vim pv.yaml
root@master:~# kubectl create -f pv.yaml
persistentvolume/pv-analytics created
root@master:~# kubectl get pv
NAME                CAPACITY   ACCESS MODES   RECLAIM POLICY   STATUS   CLAIM   STORAGECLASS   REAS
ON  AGE
pv-analytics        100Mi      RWX            Retain           Available
8s
root@master:~#
```

Read More: [K8s Persistent Storage](#)

Q5) Taint the worker node to be Unschedulable. Once done, create a pod called dev-redis, image redis:alpine to ensure workloads are not scheduled to this worker node. Finally, create a new pod called prod-redis and image redis:alpine with toleration to be scheduled on node01.

key:env_type, value:production, operator: Equal and effect:NoSchedule

Ans.

```
$ kubectl get nodes
$ kubectl taint node node01 env_type=production:NoSchedule
$ kubectl describe nodes node01 | grep -i taint
$ kubectl run dev-redis --image=redis:alpine --dry-run=client -o yaml >
pod-redis.yaml
$ vi prod-redis.yaml
apiVersion: v1
kind: Pod
metadata:
  name: prod-redis
spec:
  containers:
  - name: prod-redis
    image: redis:alpine
  tolerations:
  - effect: NoSchedule
    key: env_type
    operator: Equal
    value: production
$ kubectl create -f prod-redis.yaml
```

Read More: [Scheduling in K8s](#)

Q6) Set the node named worker node as unavailable and reschedule all the pods running on it. (Drain node)

Ans.

```
$ kubectl drain node <worker node> --ignore-daemonsets
```

Q7) Create a Pod called non-root-pod , image: redis:alpine

runAsUser: 1000

fsGroup: 2000

Ans.

```
$ vim non-root-pod.yaml
$ kubectl create -f non-root-pod.yaml
apiVersion: v1
kind: Pod
metadata:
  name: non-root-pod
spec:
  securityContext:
    runAsUser: 1000
    fsGroup: 2000
  containers:
  - name: non-root-pod
```

Read More: [K8s Pods For Beginners](#)

Q8) Create a NetworkPolicy which denies all ingress traffic

Ans.

```
$ vim policy.yaml
apiVersion: networking.k8s.io/v1
kind: NetworkPolicy
metadata:
  name: default-deny
spec:
  podSelector: {}
  policyTypes:
  - Ingress
$ kubectl create -f policy.yaml
```

Read More: [K8s Network Policy](#)

Conclusion

Kubernetes is the leading technology, and companies always look for skilled employees. To help you crack the CKA exam and secure a job, we put some effort and listed some Sample Exam Questions.